CS57300: Assignment 4

Due date: November 8 (Friday), 11:59 pm (submit via BlackBoard)

Comparing Methods for Speed Dating Classification

In this programming assignment, you will be asked to implement **Decision Trees**, **Bagged Trees** and **Random Forests** for the classification task that you explored in Assignments 2 and 3, and then compare the performance of different classifiers. Bagged Trees and Random Forest will require the training of multiple decision trees. Generating learning curves by tuning different parameters for the above mentioned three implementations will require even more time and effort. With these in mind, we highly recommend you to start this assignment early!

Similar as that in Assignment 2 and 3, you must design and implement your **own versions** of the algorithm in Python for this assignment. **DO NOT** use any publicly available code including libraries such as **sklearn**. Your code will be checked against public implementations. In addition, we will not provide separate testing data to you. You are asked to design your own tests to ensure that your code runs correctly and meets the specifications below.

Note: You may use the **pandas, numpy, scipy** libraries for data processing purposes. The only restriction is that you have to write your **own version** of data mining algorithms; you can not use any built-in functions for your algorithm. This is a general rule for this assignment and all the upcoming ones as well. As before, you should submit your typed assignment report as a pdf along with your source code file.

In the following sections, we specify a number of steps you are asked to complete for this assignment. **Note that all results in sample outputs are fictitious and for representation only**.

1 Preprocessing

Consider the data file **dating-full.csv** that you used in Assignment 2. For this assignment, we will only consider the first 6500 speed dating events in this file. That is, you can discard the last 244 lines of the file. Write a Python script named **preprocess-assg4.py** that reads the first 6500 speed dating events in **dating-full.csv** as input and performs the following operations.

(i) For simplicity, drop the columns **race**, **race_o** and **field**.

(ii) For the categorical attribute gender, apply label encoding, as in Assignment 2.

(iii) Repeat the preprocessing steps 1(iv) that you did in Assignment 2. (You can reuse the code there and you are not required to print any outputs.)

(iv) Discretize all the continuous-valued columns (the columns mentioned as continuous valued columns in Assignment 2) using 2 bins of equal widths (You should use the cut function in **pandas** with number of bins as 2 and labels = [0, 1] to do so), so that all the features values become binary (0/1).

(v) Use the **sample** function from **pandas** with the parameters initialized as **random_state = 47, frac = 0.2** to take a random 20% sample from the entire dataset. This sample will serve as your test dataset, which you should output in **testSet.csv**; the rest will be your training dataset, which you should output in **trainingSet.csv**. (Note: The use of the random_state

will ensure all students have the same training and test datasets; incorrect or no initialization of this parameter will lead to non-reproducible results).

2   Implement Decision Trees, Bagging and Random Forests (10 points)

Please put your code for this question in a file called **trees.py**. This script should take three arguments as input:

1. *trainingDataFilename*: the set of data that will be used to train your algorithms (e.g., **trainingSet.csv**).

2. *testDataFilename*: the set of data that will be used to test your algorithms (e.g., **testSet.csv**).

3. *modelIdx*: an integer to specify the model to use for classification (DT = 1, BT = 2, RF = 3, where DT refers to decision trees, BT refers to bagging, and RF refers to random forests).

(i) Write a function named ***decisionTree(trainingSet, testSet)*** that takes the training dataset and the testing dataset as input parameters. The purpose of this function is to train a decision tree classifier using the data in the training dataset, and then test the classifier's performance on the testing dataset.

   Use Gini-gain as a score function. Grow trees using a depth limit of 8 and an example limit of 50 (i.e., stop growing when either the depth of the tree reaches 8 or the number of examples in a node is smaller than 50). When you reach either of the stopping criteria, consider that node to be the leaf node that is able to make prediction.

   Do not confuse between depth of a node and level of a node. The depth of the root node in a tree is 0, whereas the level of the root is 1. **You will be using the notion of depth here.**

(ii) Write a function named ***bagging(trainingSet, testSet)*** that takes the training dataset and the testing dataset as input parameters. The purpose of this function is to train a bagged decision tree classifier using the data in the training dataset, and then test the classifier's performance on the testing dataset.

   Learn 30 trees with the stopping criterion in (i). Use sampling with replacement to construct pseudosamples (i.e., bootstrapped sample of the training data).

(iii) Write a function named ***randomForests(trainingSet, testSet)*** that takes the training dataset and the testing dataset as input parameters. The purpose of this function is to train a random forests classifier using the data in the training dataset, and then test the classifier's performance on the testing dataset.

   Learn 30 trees with the stopping criterion in (i). Use sampling with replacement to construct pseudosamples. Use $\sqrt{p}$ to downsample the features at each node of the tree (where p is the total number of features), that is, at a particular node, sample $\sqrt{p}$ features, then exclude any of those features used in the path from root to that node.

The sample inputs and outputs we expect to see are as follows (the numbers are fictitious):

```
$python trees.py trainingSet.csv testSet.csv 1
Training Accuracy DT: 0.71
```

```
Testing Accuracy DT: 0.68

$python trees.py trainingSet.csv testSet.csv 2
Training Accuracy BT: 0.75
Testing Accuracy BT: 0.74

$python trees.py trainingSet.csv testSet.csv 3
Training Accuracy RF: 0.73
Testing Accuracy RF: 0.77
```

3   The Influence of Tree Depth on Classifier Performance (10 points)

Please follow the procedure below to assess whether the depth of the tree affects classifier performance. Put your code for this question in a file called **cv_depth.py**.

Use the **sample** function from **pandas** with the parameters initialized as **random_state = 18, frac = 1** to shuffle the training data (i.e., data in **trainingSet.csv**). Then, obtain a 50% sample of the above shuffled training data using **random state=32**.

Perform 10-fold cross validation on this sample of training data just created above (consider the first 10% lines of the sampled data as your first fold, the second 10% lines of the sampled data as your second fold, and so on; notice you are asked to conduct cross validation on a sample of the training data to reduce the time cost of this assignment). Conduct the cross validation for each of the three models—decision tree, bagged trees, and random forests—where depth limit of the trees in each model is set to be $d \in [3, 5, 7, 9]$. The example limit of the trees is fixed at 50. Learn 30 trees for the ensemble models.

(a) Plot the average accuracy for 10-fold cross validation on y-axis, and depth limit of tree on x-axis. Include error bars that indicate ± 1 standard error (use the formula used in Assignment 3). Please include the curves for the three models in one figure.

(b) Formulate a hypothesis about the performance difference (of any two models of your choice) you observe as the depth limit of trees change. Discuss whether the observed data support the hypothesis or not (i.e., are the observed differences significant?). Use significance level (alpha value) 0.05.

4   Compare Performance of Different Models (10 points)

Please follow the procedure below to assess whether ensembles improve performance. Put your code for this question in a file called **cv_frac.py**.

Use the **sample** function from **pandas** with the parameters initialized as **random state = 18, frac = 1** to shuffle the training data (i.e., data in **trainingSet.csv**). Conduct 10-fold incremental cross validation, as described in Assignment 3, for the three models.

Specifically, please first divide the shuffled training data into 10 folds (i.e., the first 10% lines is your first fold, the second 10% lines is your second fold, and so on). Then, use fractions $t\_frac \in [0.05, 0.075, 0.1, 0.15, 0.2]$ with **random state=32** to obtain $train\_set$ in each iteration of cross validation. The depth limit of tree is 8, and example limit is 50. Learn 30 trees for the ensemble models.

(a) Plot the learning curves for the three models (in the same plot), with the average accuracy of the 10 trials on y-axis, and the training fraction on x-axis. Include error bars to indicate ± 1 standard error (see descriptions in Assignment 3, Q3(ii)(b) for how to compute standard error).

(b) Formulate a hypothesis about the performance difference you observe between the decision tree and any one of the ensemble methods. Discuss whether the observed data support the hypothesis or not (i.e., are the observed differences significant?). Use significance level (alpha value) 0.05.

5 The Influence of Number of Trees on Classifier Performance (10 points)

Please following the procedure below to assess whether the number of trees affects performance. Put your code for this question in a file called **cv_numtrees.py**.

Use the **sample** function from **pandas** with the parameters initialized as **random_state = 18, frac = 1** to shuffle the training data (i.e., data in **trainingSet.csv**). Then, obtain a 50% sample of the above shuffled training data using **random state=32**.

Perform 10-fold cross validation on this sample of training data (consider the first 10% lines of the sampled data as your first fold, the second 10% lines of the sampled data as your second fold, and so on; notice you are asked to conduct cross validation on a sample of the training data to reduce the time cost of this assignment). Conduct the cross validation for each of the two ensemble methods—bagged trees and random forests—where the number of trees in each model is set to be $t \in [10, 20, 40, 50]$. The depth limit of tree is 8, and example limit is 50.

(a) Plot the average accuracy for 10-fold cross validation on y-axis, and number of trees on x-axis. Include error bars that indicate ± 1 standard error. Please include the curves for the two models in one figure.

(b) Formulate a hypothesis about the performance difference you observe between the two ensemble models. Discuss how the observed data support the hypothesis (i.e., are the observed differences significant?). Use significance level (alpha value) 0.05.

Bonus question (5 points)

Implement a suitable model of your choice that has not been included in Assignments 2, 3, or 4, (e.g., boosted decision trees, neural networks, etc.), along with the optimal set of hyper-parameters, that gives highest possible accuracy on the testing dataset (testSet.csv) (recall that you should not touch the testing dataset until you are satisfied with your model). Report your tuning procedure, the hyper-parameters you end up with, your model selection procedure, training and testing procedures, and the level of accuracy you get on the testing dataset in the report. Also include the coding files and mention their names in the report. **Note that you have to implement the complete model without using any available softwares such as Weka, or libraries like sklearn.**

Submission Instructions:

Please submit your assignment through Blackboard:

1. Include in your report which version of Python you are using.

2. Make sure you include in your report all the output / results you get from running your code for all sub-questions. You may include screen shots to show them.

3. Make a directory named *yourFirstName_yourLastName_HW4* and copy all of your files to this directory.

4. **DO NOT** put the datasets into your directory.

5. Make sure you compress your directory into a **zip folder** with the same name as described above, and then upload your zip folder to Blackboard.

   Keep in mind that old submissions are overwritten with new ones whenever you re-upload.

 Your submission should include the following files:

1. The source code in python.

2. Your evaluation & analysis in .pdf format. Note that your analysis should include visualization plots as well as a discussion of results, as described in details in the questions above.

3. A README file containing your name, instructions to run your code and anything you would like us to know about your program (like errors, special conditions, etc).